

# 13 Ant Colony Optimisation

# Ant Colony Optimisation

- introduced by Dorigo in 1992
- inspired by the foraging behaviour of ants
- because ants are interesting!
  - they solve complex tasks by simple local means
  - ant productivity is better than the sum of their single activities
  - ants are 'grand masters' in search and exploitation
- for example, they are able to find the shortest path from food to nest
  - but how do they do it?

# Ant Colony Optimisation

信息素

- to find the shortest path they use pheromone
- ants deposit pheromone as they move along a path
- this is used by other ants to help them decide which route to take
  - who consequently lay even more pheromone along the same path
  - a positive feedback mechanism
- so they alter the local environment in order to indirectly communicate with each other
  - a process known as stigmergy - ‘stimulation by work’
- and their combined actions lead towards optimal solutions
- this leads to robustness and adaptability


Stigmergy (斯蒂格默基) 是一种社会学和生物学概念，用于描述在没有任何集中式控制或通信的情况下，一组个体如何通过间接相互作用来协同工作。在中文中，通常将其翻译为“诱导性协作”、“诱导协同”或“诱导合作”。

# Ant Foraging Example

- suppose there are two paths from nest to food
- initially each has no pheromone laid on it
  - so each is equally likely to be chosen
- two ants set off from the nest, one down each path...


# Ant Foraging Example

two ants set off from the nest, one down each path...




# Ant Foraging Example

...the ant on the shorter path can make more journeys...




# Ant Foraging Example

...so over time it lays more pheromone on its path than the other ant does...



# Ant Foraging Example


...so the next ant out of the nest is more likely to follow the shorter route...






# Ant Foraging Example

...and as more ants leave the nest, more are likely to use the path with the higher pheromone level, further reinforcing it....



# Ant Foraging Example

...until the shorter path is almost always used.



# Evaporation

- so, in the real world, ants initially wander randomly, and upon finding food return to their colony
  - all the time laying down pheromone trails
- other ants are likely to follow an existing trail, returning and reinforcing it if they eventually find food
- over time, however, the pheromone trail starts to evaporate
  - which reduces its attractive strength
- the more time it takes for an ant to travel down the path and back again, the more time the pheromones have to evaporate

# Evaporation Helps Avoid Local Optima

- a short path gets marched over faster than a long path
- so the pheromone density remains high
  - it's laid at least as fast as it can evaporate
- pheromone evaporation is crucial for avoiding convergence to a locally optimal solution
- without evaporation the paths chosen by the first ants would tend to be excessively attractive to the following ones
- so exploration of the solution space would be constrained

# ACO Algorithm

- a population based metaheuristic
- used to find approximate solutions to difficult optimisation problems
- the optimisation problem is encoded as the environment
  - usually a graph, with weighted edges
- a population of ants is created, usually divided into generations

# ACO Algorithm

- each generation of ants traverse the graph and builds their own solutions
  - find their own paths
- each ant lays pheromone on edges proportionate to the quality of its solution
- successive generations converge towards optimal paths
  - helping to build an optimal solution

# ACO Algorithm

- the basic algorithm begins with initialisation:
  - set up nodes
  - create edges between nodes, and for each edge:
    - set edge weighting
    - set pheromone level
- these three procedures are then performed each generation:
  - Generate Ant Solutions
  - Perform Daemon Actions (optional)
  - Update Pheromone Levels

# ACO: Generate Ant Solutions

- this manages the population (colony) of ants
- each generation a set of ants are created and placed at nodes
  - perhaps randomly
- each ant travels through the graph until its task is complete
- an ant needs to perform edge selection to decide which node to visit next
  - edge selection is stochastic and based on:
    - pheromone levels
    - heuristic information (such as edge-weighting)
- each ant's solution is evaluated



# ACO: Edge Selection Formula

- selecting which edge to traverse next is probabilistic
- the best known rule is ant system (AS), created by Dorigo et al in 1996:

$$\text{prob}(\text{choose edge } i) = (\tau_i^\alpha \cdot \eta_i^\beta) / \sum_n (\tau_i^\alpha \cdot \eta_i^\beta)$$

where:


$n$  = number of edges

$\tau_i$  = level of pheromone on edge  $i$  'trail level'

$\eta_i$  =  $1/(\text{length of edge } i)$  'attractiveness'


$\alpha, \beta$  are positive real parameters that determine the relative importance of pheromone versus heuristic information (edge length)

# ACO: Edge Selection Example



- ant at node  $i$
- 3 edges lead to nodes **not yet visited**
- pheromone level shown in purple
- edge length shown in blue

# ACO: Edge Selection Example



using weightings:

$$\alpha = 1$$

$$\beta = 2$$

edge	$\tau$	$\tau^\alpha$	$\eta = 1 \div \text{length}$	$\eta^\beta$	$\tau^\alpha \cdot \eta^\beta$	prob = $(\tau^\alpha \cdot \eta^\beta) \div \Sigma$
i to j	3	3	0.2	0.04	0.12	$0.12 \div 5.37 = \mathbf{0.022}$
i to k	5	5	0.5	0.25	1.25	$0.25 \div 5.37 = \mathbf{0.233}$
i to l	4	4	1	1	4	$4 \div 5.37 = \mathbf{0.745}$
					$\Sigma=5.37$	1.00

# ACO: Perform Daemon Actions

- these optionally occur after individual ant solutions have been constructed, and before pheromone levels are updated
- they can perform actions that are problem specific or centralised
  - and so cannot be performed by individual ants
- for example, applying local search of the solutions to select which solutions are allowed to update pheromone levels

# ACO: Update Pheromone Levels

- once all the ants in a generation have completed their tour, the pheromone levels on the edges are updated
- each edge's pheromone level is initially decreased by a certain percentage
  - evaporation
- each edge then receives an amount of additional pheromone
- the amount laid is proportional to the quality of the solutions to which it belongs
- this procedure is then repeated for each generation, until a termination criterion is satisfied

# ACO: Update Pheromone Formula

- the pheromone update formula looks like this:

$$\tau_i \leftarrow (1-\rho) \cdot \tau_i + \rho \cdot \sum (\Delta\tau_i \text{ for ant } k)$$

where:

$\tau_i$  = pheromone level on edge  $i$

$\rho$  = pheromone evaporation coefficient

$\Delta\tau_i$  = amount of pheromone deposited by an ant in this generation:

$Q/L$  if the ant uses the edge in its tour


0 otherwise

where  $L$  is the tour length

$Q$  is a constant


- in words:
  - perform evaporation on the current pheromone level
  - add the new pheromone laid by each ant

# ACO: Update Pheromone Levels



- ant has completed its tour
- for the **edges travelled:**
- pheromone level shown in **purple**
- edge length shown in **blue**

# ACO: Update Pheromone Levels



using :

evaporation coefficient  $\rho = 0.1$

pheromone per ant  $Q = 10$

first evaporate pheromone on all edges in the graph

edge	initial $\tau$	after evap $\tau$
1 to 2	6 - 0.6	5.4
2 to 4	2 - 0.2	1.8
4 to 6	1 - 0.1	0.9
6 to 5	7 - 0.7	6.3
5 to 3	5 - 0.5	4.5
3 to 1	1 - 0.1	0.9

...and for all other edges (9 more)

then add pheromone for each ant's tour

length	$\Delta\tau=Q/L$	updated $\tau$
1	0.43	5.83
3	0.43	2.23
5	0.43	1.33
3	0.43	6.73
7	0.43	4.93
4	0.43	1.33
L = 23	(repeat for each ant)	

10/23



# ACO Example: Travelling Salesman Problem

- in the traveling salesman problem (TSP) a set of locations (cities) and the distances between them are given
- the problem consists of finding a closed tour of minimal length that visits each city once and only once
- to apply ACO to the TSP we represent:
  - cities as nodes
  - distances between cities as the weightings (lengths) of edges between those nodes
- this graph is known as the construction graph


# ACO for TSP

- each generation, a set of ants can be placed on random nodes in the graph
- each ant then tours the graph, visiting each city once, and recording its tour
- it selects edges by looking at:
  - whether or not the edge leads to a node already visited (valid or invalid)
  - for valid edges it probabilistically chooses which edge to traverse based on:
    - the amount of pheromone on each edge
    - the length of each edge (longer edges are less likely to be chosen)


# ACO for TSP

- at the end of a generation each ant lays pheromone on the edges it traversed during its tour
- the amount of pheromone laid per edge is inversely proportional to the length of the tour
  - so shorter tours lay more pheromone per edge
  - while longer tours lay less pheromone per edge
- this information encourages ants in future generations to select edges that formed part of shorter tours
- and hence the ants collectively contribute towards finding an optimal solution: the shortest tour of the graph


# ACO for TSP: Tour de France




**1**



**2**



**3**

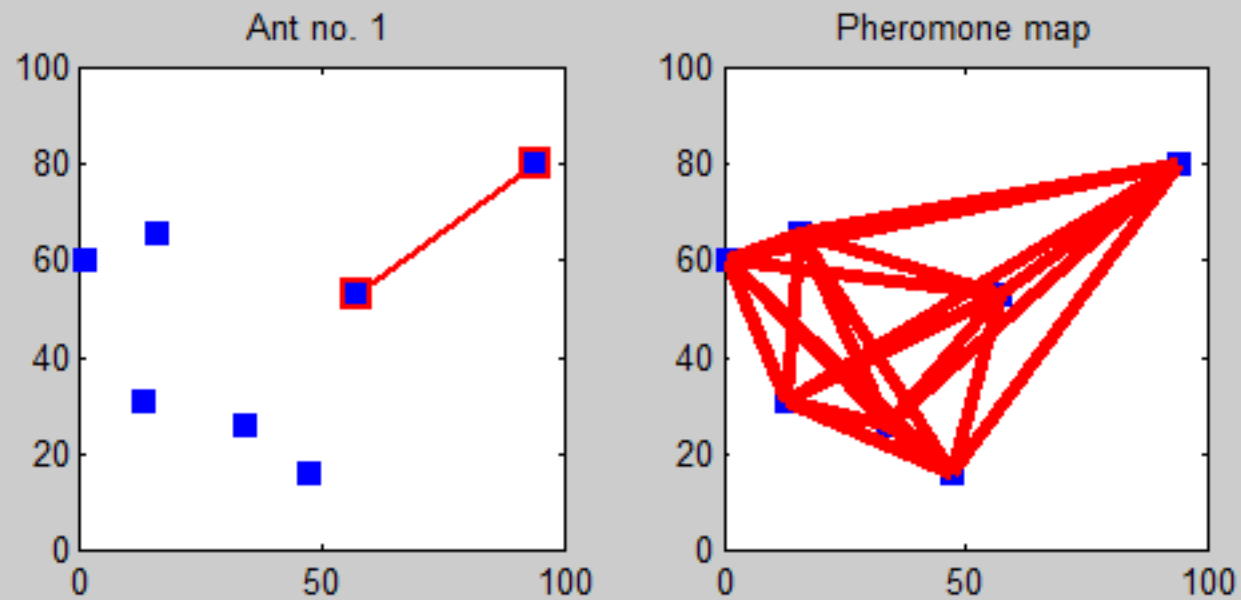


**4**

# Example Run: ACO for TSP




# Example Run: ACO for TSP




start of run

# Example Run: ACO for TSP



middle of run

# Example Run: ACO for TSP



end of run



# Main Types of ACO

- the three most common types of ACO are:
- Ant System (AS)
  - introduced by Dorigo et al, 1996
- Ant Colony System (ACS)
  - Dorigo et al, 1997
- MAX-MIN ant system (MMAS)
  - Stützle & Hoos, 2000

# Ant System (AS)

- the original ACO algorithm
- main characteristic:
  - pheromone values are updated by *all* ants that have completed the tour
- edge selection and pheromone update formulae are those seen in the previous slides

# Ant Colony System (ACS)

- a major improvement over Ant System
- ants in ACS use a different edge selection rule:
- the **pseudorandom proportional rule**:
- the probability for an ant to move from city  $x$  to city  $y$  depends on:
  - a random variable  $q$ , uniformly distributed over  $[0, 1]$
  - a parameter  $q_0$ :
    - if  $q \leq q_0$  then the edge that maximises  $\tau_{xy} \cdot \eta_{xy}^\beta$  is chosen
    - otherwise the same formula as AS is used
- this rule **favours exploitation of the pheromone information**

# ACS: Local Pheromone Update

- the exploitation provided by the edge selection rule is counterbalanced by the **local pheromone update**:
- *during a tour*, after each edge is traversed an ant updates that edge's pheromone level:

$$\tau_i \leftarrow (1-\rho) \cdot \tau_i + \rho \cdot \tau_{init}$$

- where:
  - $\rho \in (0, 1]$  is the pheromone decay coefficient
  - $\tau_{init}$  is the initial value of the pheromone

# ACS: Local Pheromone Update

- the main aim of the local update is to diversify the search performed by subsequent ants during one iteration
- decreasing the pheromone concentration on the edges as they are traversed during a generation encourages subsequent ants to choose other edges
  - and hence to produce different solutions
  - so it's less likely that several ants produce identical solutions during a single generation

# ACS: Offline Pheromone Update

- at the end of a generation a pheromone update is performed
- this greatly differs from Ant System
  - in AS all ants that traversed an edge update the pheromone on it
  - however in ACS this offline update is only performed by the best ant of this generation

# MAX-MIN Ant System (MMAS)

- MAX-MIN ant system (MMAS) is another improvement, proposed by Stützle and Hoos (2000), over the original ant system idea
- MMAS differs from AS in that:
  - only the best ant adds pheromone trails
  - the minimum and maximum values of the pheromone are explicitly limited ( $\tau_{\max}$  and  $\tau_{\min}$ )
    - these values are often experimentally chosen
- in AS and ACS the max and min values are limited implicitly:
  - the value of the limits is a result of the algorithm working
  - rather than a value set explicitly by the algorithm designer

# Benefits of ACO

- ACO algorithms can produce near-optimal solutions to the traveling salesman problem
- they have an advantage over simulated annealing and genetic algorithm approaches when the graph can change dynamically
- because the ant colony algorithm can be run continuously and adapt to changes in real time
- this is particularly valuable to network routing and urban transportation systems



# When is it worth doing ACO?

- when a solution is easy to find
- but a *good* solution is hard to find
  - (why?)
- when you potentially want to find several good solutions
- and note that often problems that can be represented well in ACO can be easily translated into EC problems, and vice-versa
- so performance comparisons can be made

# Reading & References

- slides based on and adapted from:
  - *Ant Colony Optimization* (Riddle, University of Auckland)
  - *Swarm Intelligence* (Mohitz et al)
- recommended reading:
  - "Swarm Intelligence", Corne et al, Handbook of Natural Computing, pp 1599-1622, Springer
    - Swarm Computing module in Brightspace (TBD)
- go and play with Boids!
  - javascript (online sim): <http://www.harmendeweerd.nl/boids/>
  - java code (github): <https://github.com/tofti/javafx-boids>